

House price prediction using Multiple Linear Regression and Keras Neural Network model

Arindam Saha, 002011701091, Production Engineering

Under the guidance of Professor Dr. Debamalya Banerjee

Contents

- Introduction
- Multiple Linear Regression Model
- Neural Network Model
- Future Aspects
- References

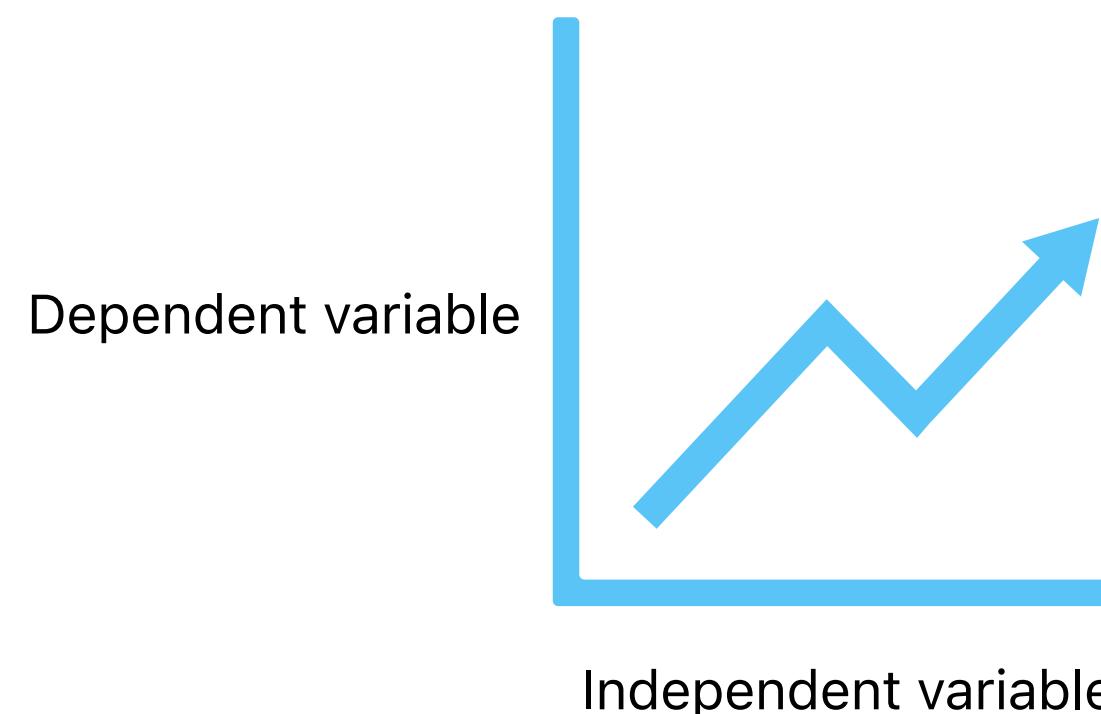
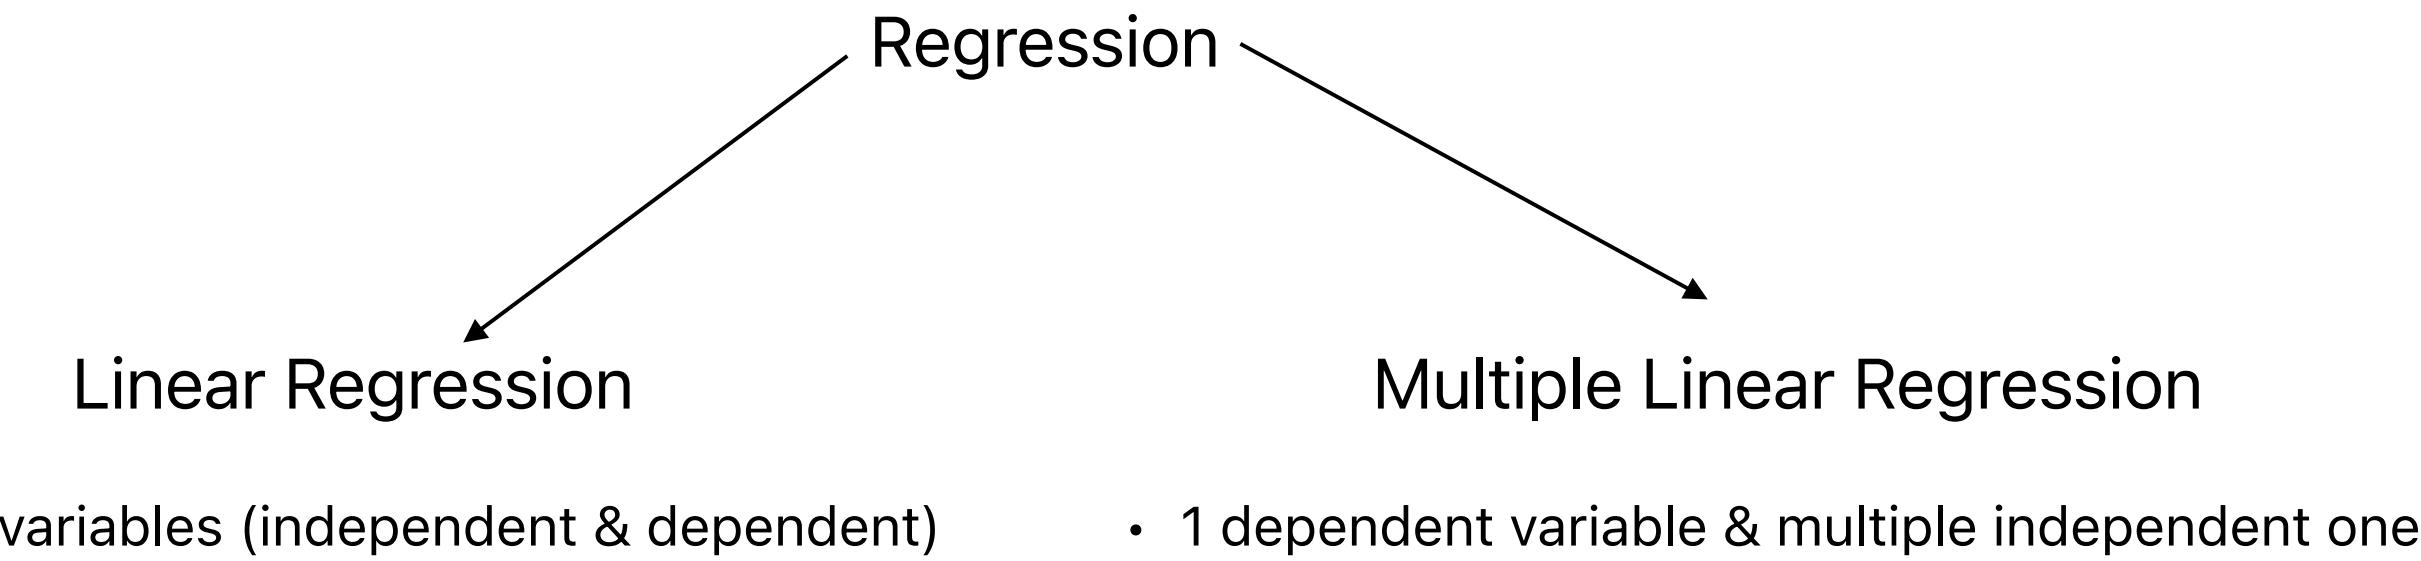
Advantages of House price prediction

- By having accurate predictions of house prices, individuals and organizations can make more informed decisions when buying or selling properties. This can lead to better investment decisions and increased profitability.
- Predictions of house prices can provide valuable insights into the real estate market, allowing individuals and organizations to better understand market trends and make more informed decisions.
- Predictions of house prices can save time by eliminating the need to manually analyze market data, allowing individuals and organizations to focus on other important tasks.
- Predictions of house prices can assist in forecasting future market trends, which can be used to make informed decisions about future investments.

Multiple Linear Regression

- Regression is the quantitative measure of relationship between dependent and independent variables.

-



Multiple Linear Regression is given by the following eq:

$$y = b_0 + b_1x_1 + b_2x_2 + \dots + b_nx_n - E$$

Assumptions

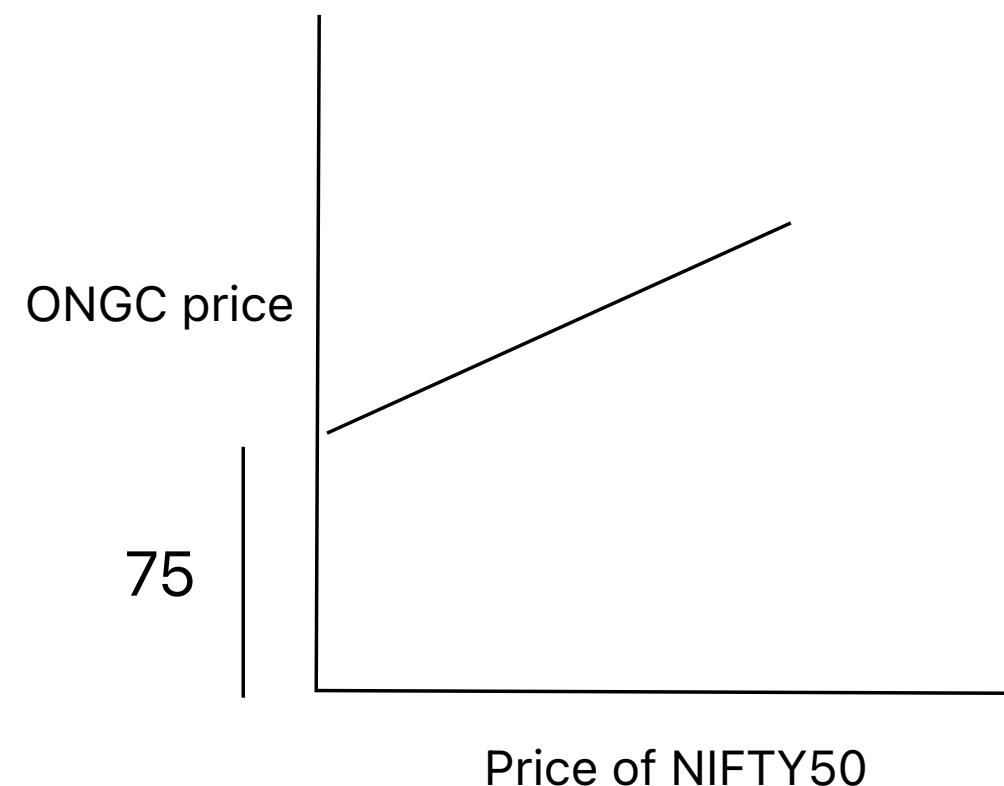
1. Linear relationship between two independent variables must be present
2. Independent variables must not be highly correlated to each other
3. y observations are selected randomly and independently from population
4. Residuals must be normally-distributed, with a mean of 0 and standard deviation of 1

For example, suppose the price of ONGC share price. It depends on multiple factors including the price of NIFTY50, price of oil, interest rate, price of future derivatives, stock price of other oil companies etc.

Suppose the equation can be given by the following -

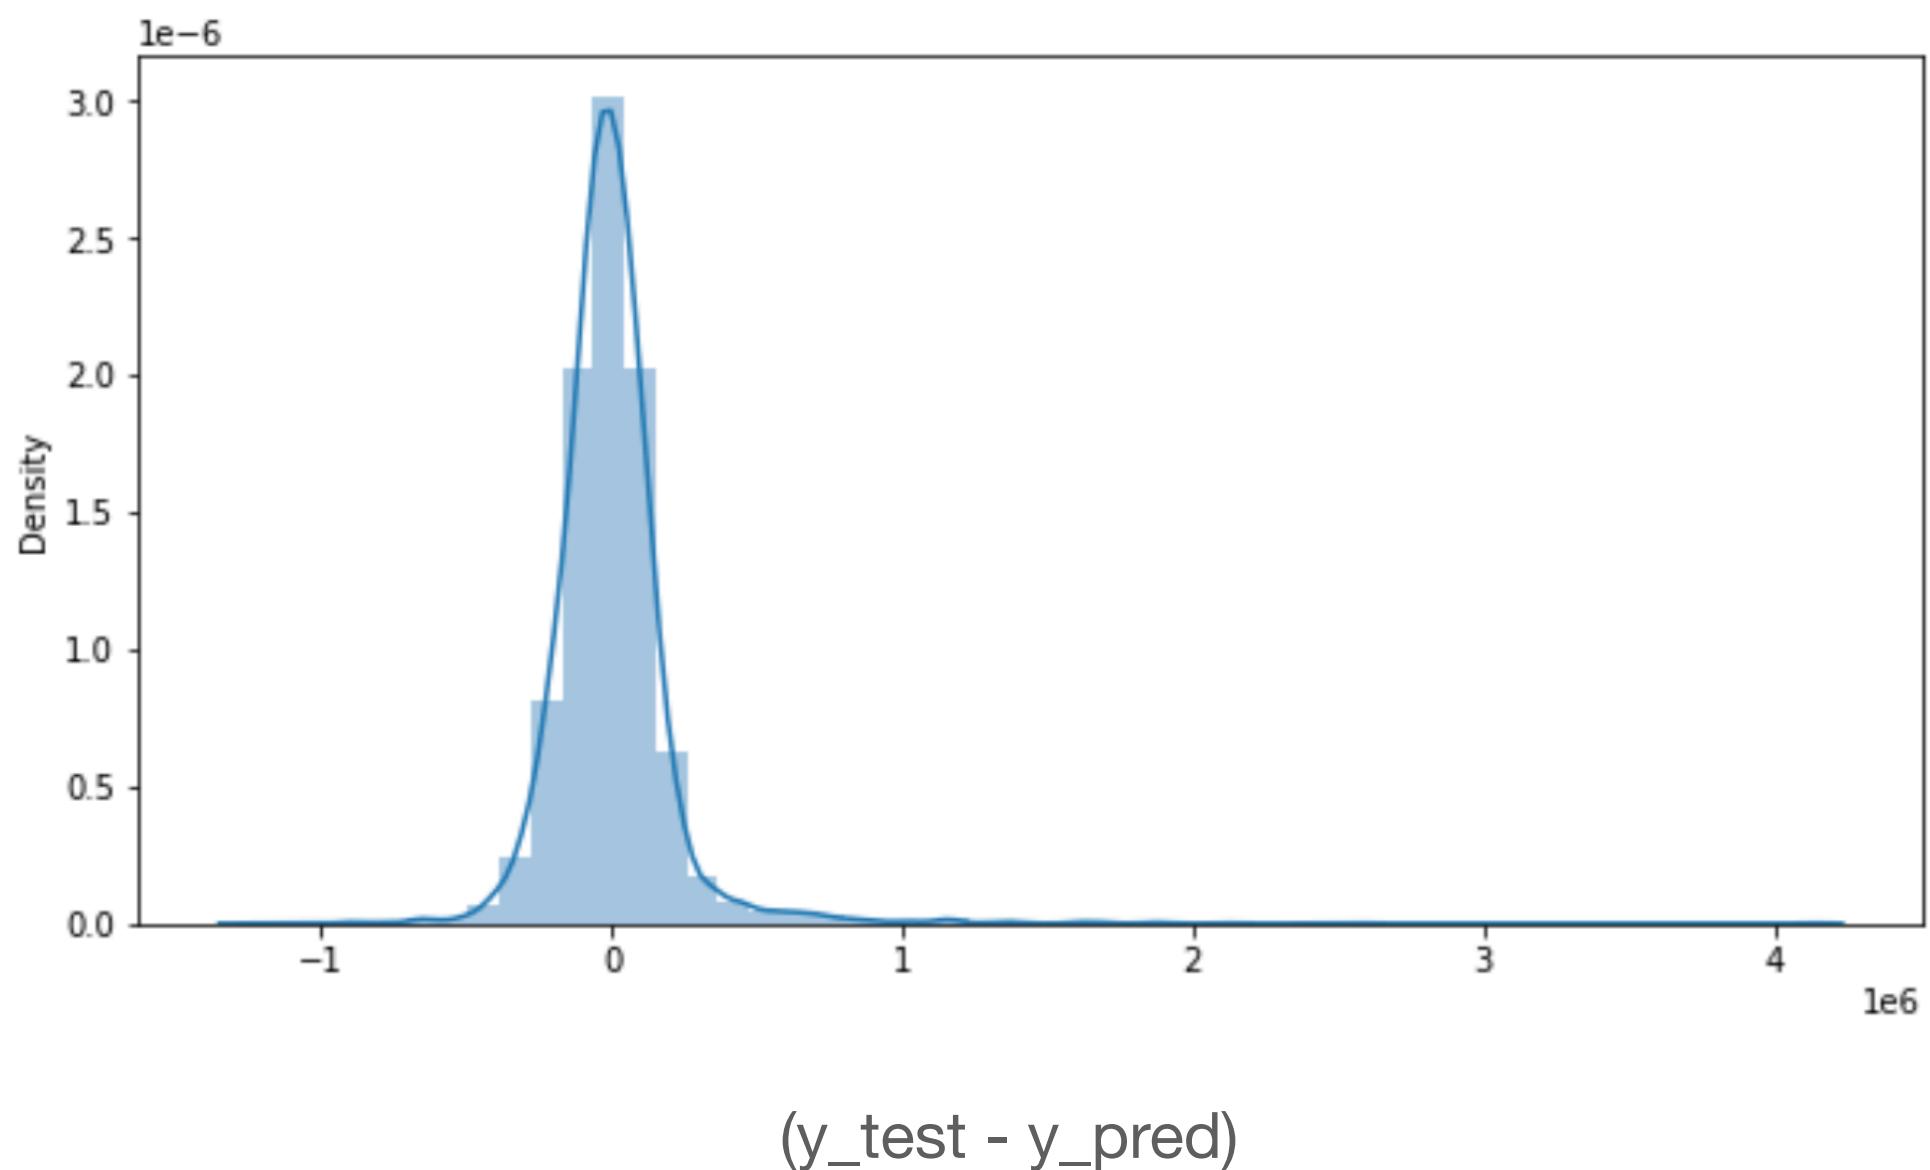
$$\text{ONGC price} = (75 - 1.5 * i + 7.8 * \text{oil price} + 3.2 * \text{NIFTY50} + 5 * \text{oil futures} - E)$$

And R-sq = 83.4%



Implementation in Python

- Get independent (x) and dependent variable (y) from dataset
- Split train & test dataset
- Feature Scaling
- Do linear regression on x & y train splits
- Find intercept & slope
- Predict test result
- Put result in a data frame
- Check accuracy and errors

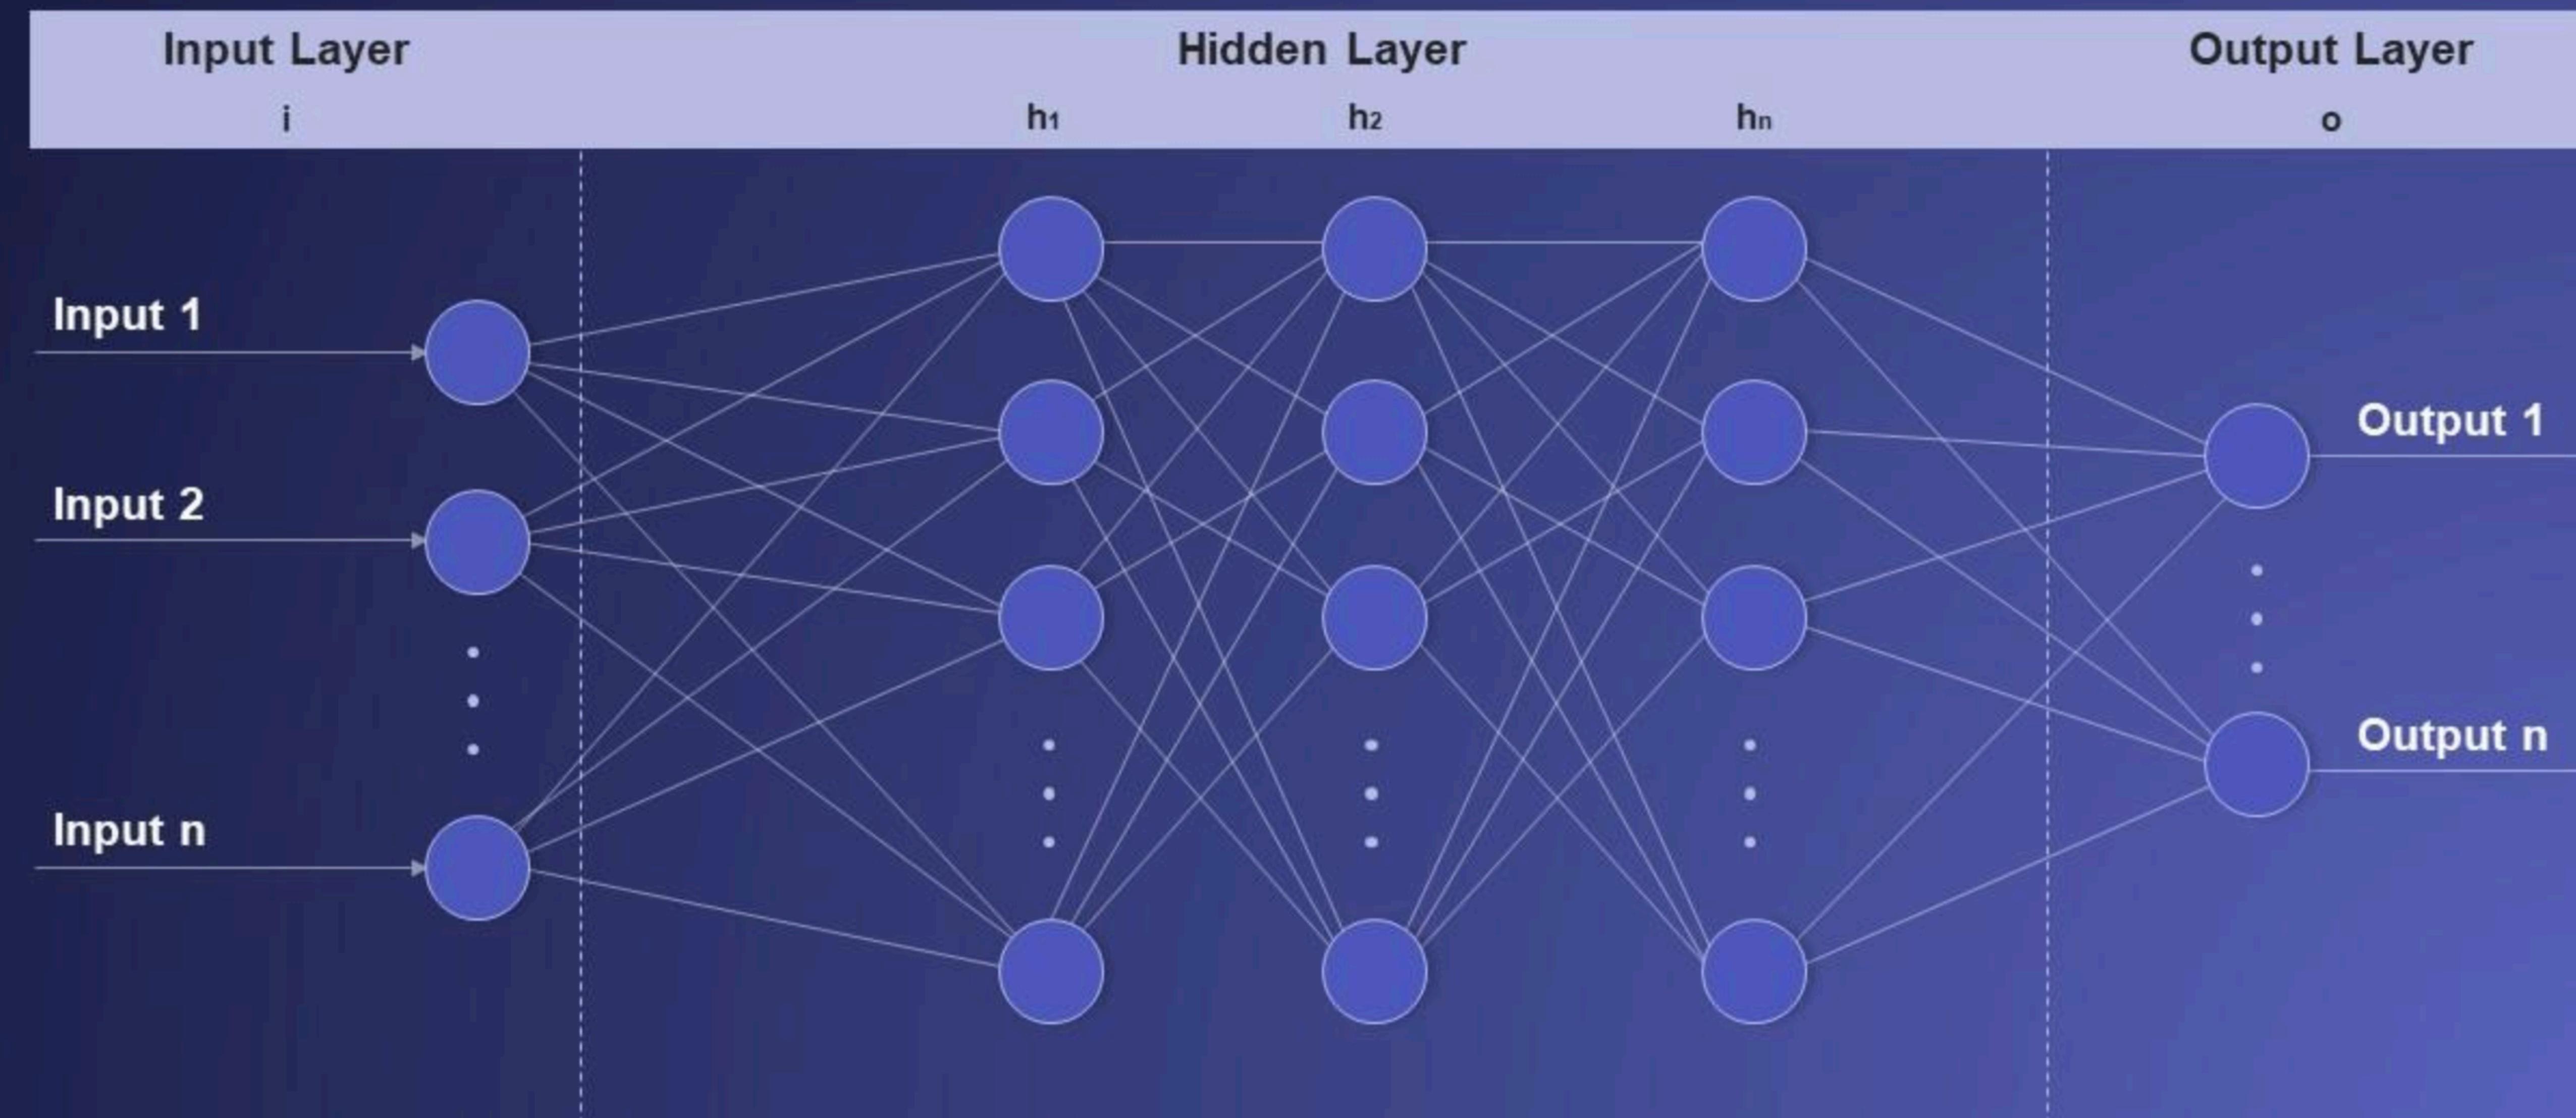


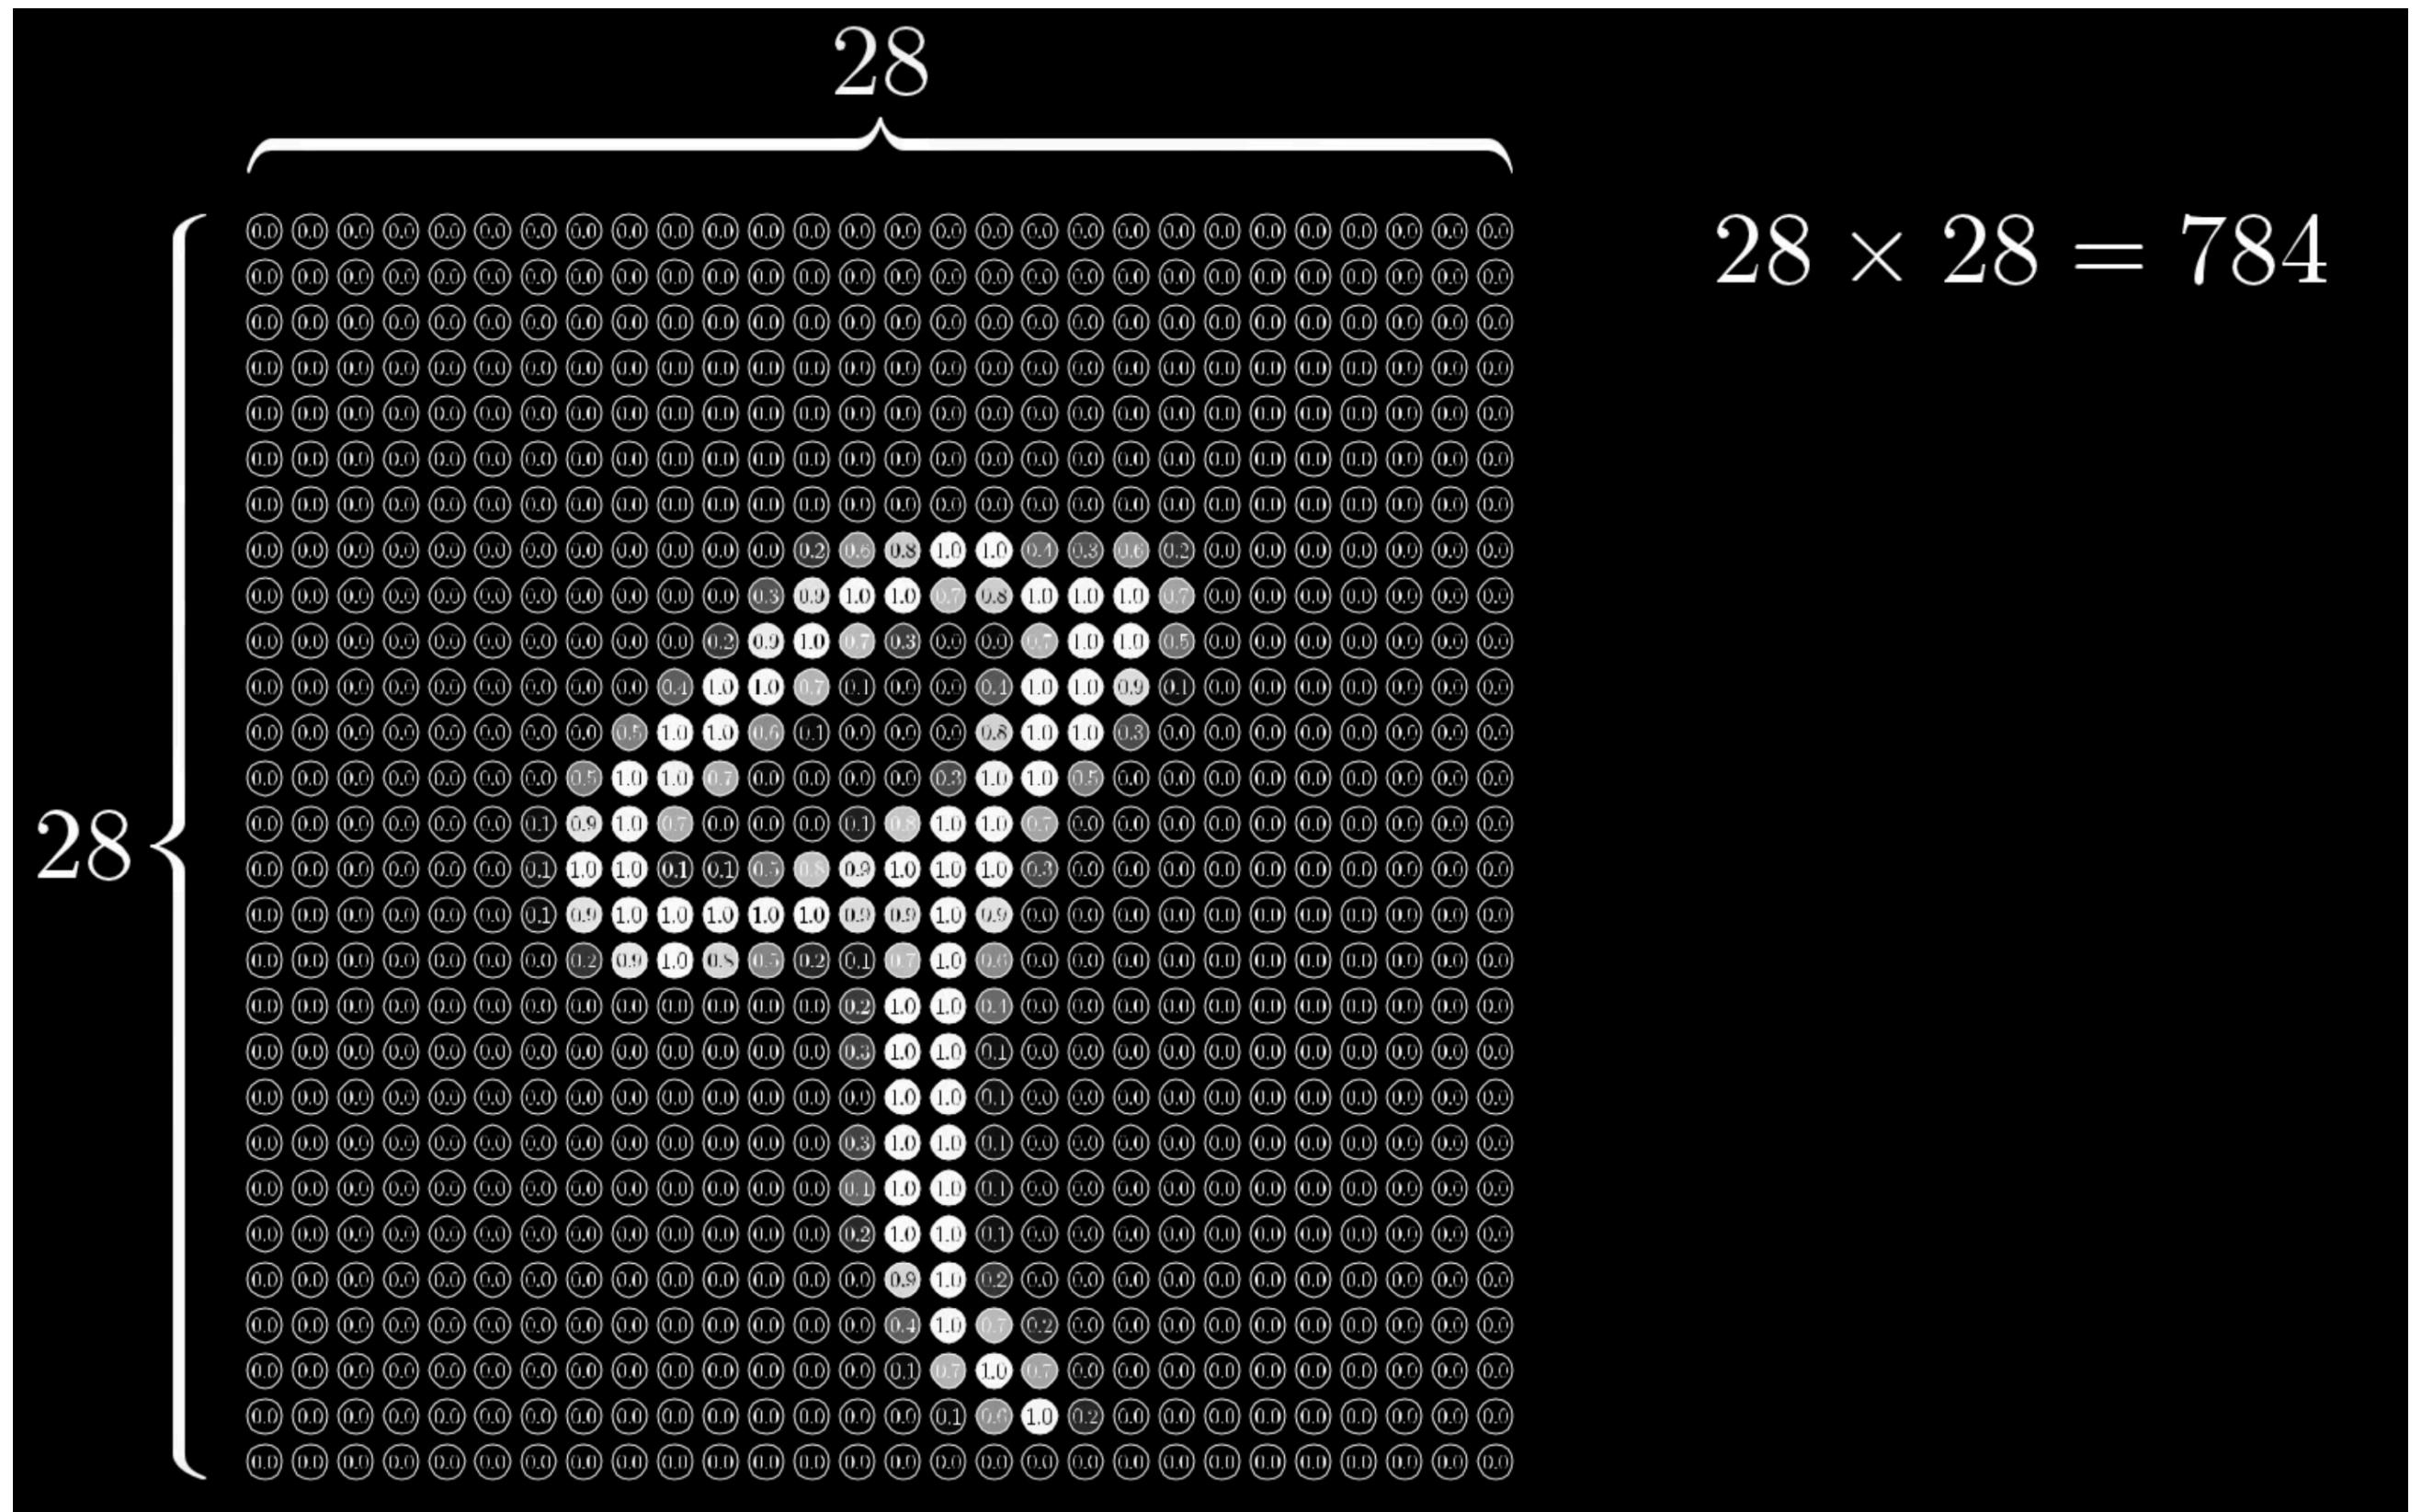
```
Price Prediction.py > ...
105
106     """### Preparing Dataset (Scaling and Train-Test Split)"""
107
108     X = Data.drop('price',axis =1).values
109     y = Data['price'].values
110
111     #splitting Train and Test
112     from sklearn.model_selection import train_test_split
113     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=101)
114
115     """## Feature Scaling"""
116
117     #standardization scaler - fit&transform on train, fit only on test
118     from sklearn.preprocessing import StandardScaler
119     s_scaler = StandardScaler()
120     X_train = s_scaler.fit_transform(X_train.astype(np.float))
121     X_test = s_scaler.transform(X_test.astype(np.float))
122
123     """# 1. Multiple Linear Regression"""
124
125     #Liner Regression
126     from sklearn.linear_model import LinearRegression
127     regressor = LinearRegression()
128     regressor.fit(X_train, y_train)
129
130     #evaluate the model (intercept and slope)
131     regressor.intercept_
132     regressor.coef_
133
134     y_predd = regressor.predict(X_test)
135     df = pd.DataFrame({'Actual': y_test, 'Predicted': y_predd})
136     df1 = df.head(10)
137     df1
138
139     fig = plt.figure(figsize=(10,5))
140     residuals = (y_test - y_predd)
141     sns.distplot(residuals)
142
143     from sklearn import metrics
144
145     print('Mean Absolute Error: {:.2f}'.format(metrics.mean_absolute_error(y_test, y_predd)))
146     print('Mean Squared Error:{:.2f}'.format(metrics.mean_squared_error(y_test, y_predd)))
147     print('Root Mean Squared Error:{:.2f}'.format(np.sqrt(metrics.mean_squared_error(y_test, y_predd))))
148     print('Variance score is: {:.2f}'.format(metrics.explained_variance_score(y_test,y_predd)))
149
150     # we are off about 20% (comparing mean absolut error and mean of price)
151     Data['price'].mean()
```

Neural Network using Keras

Architecture of artificial neural network

This slide demonstrates the architecture of Artificial Neural Network, which includes three layers such as input, hidden, and output.



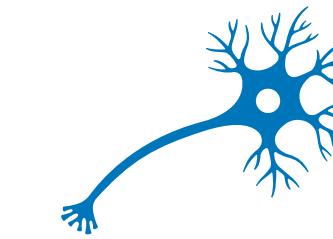


Text Recognition System

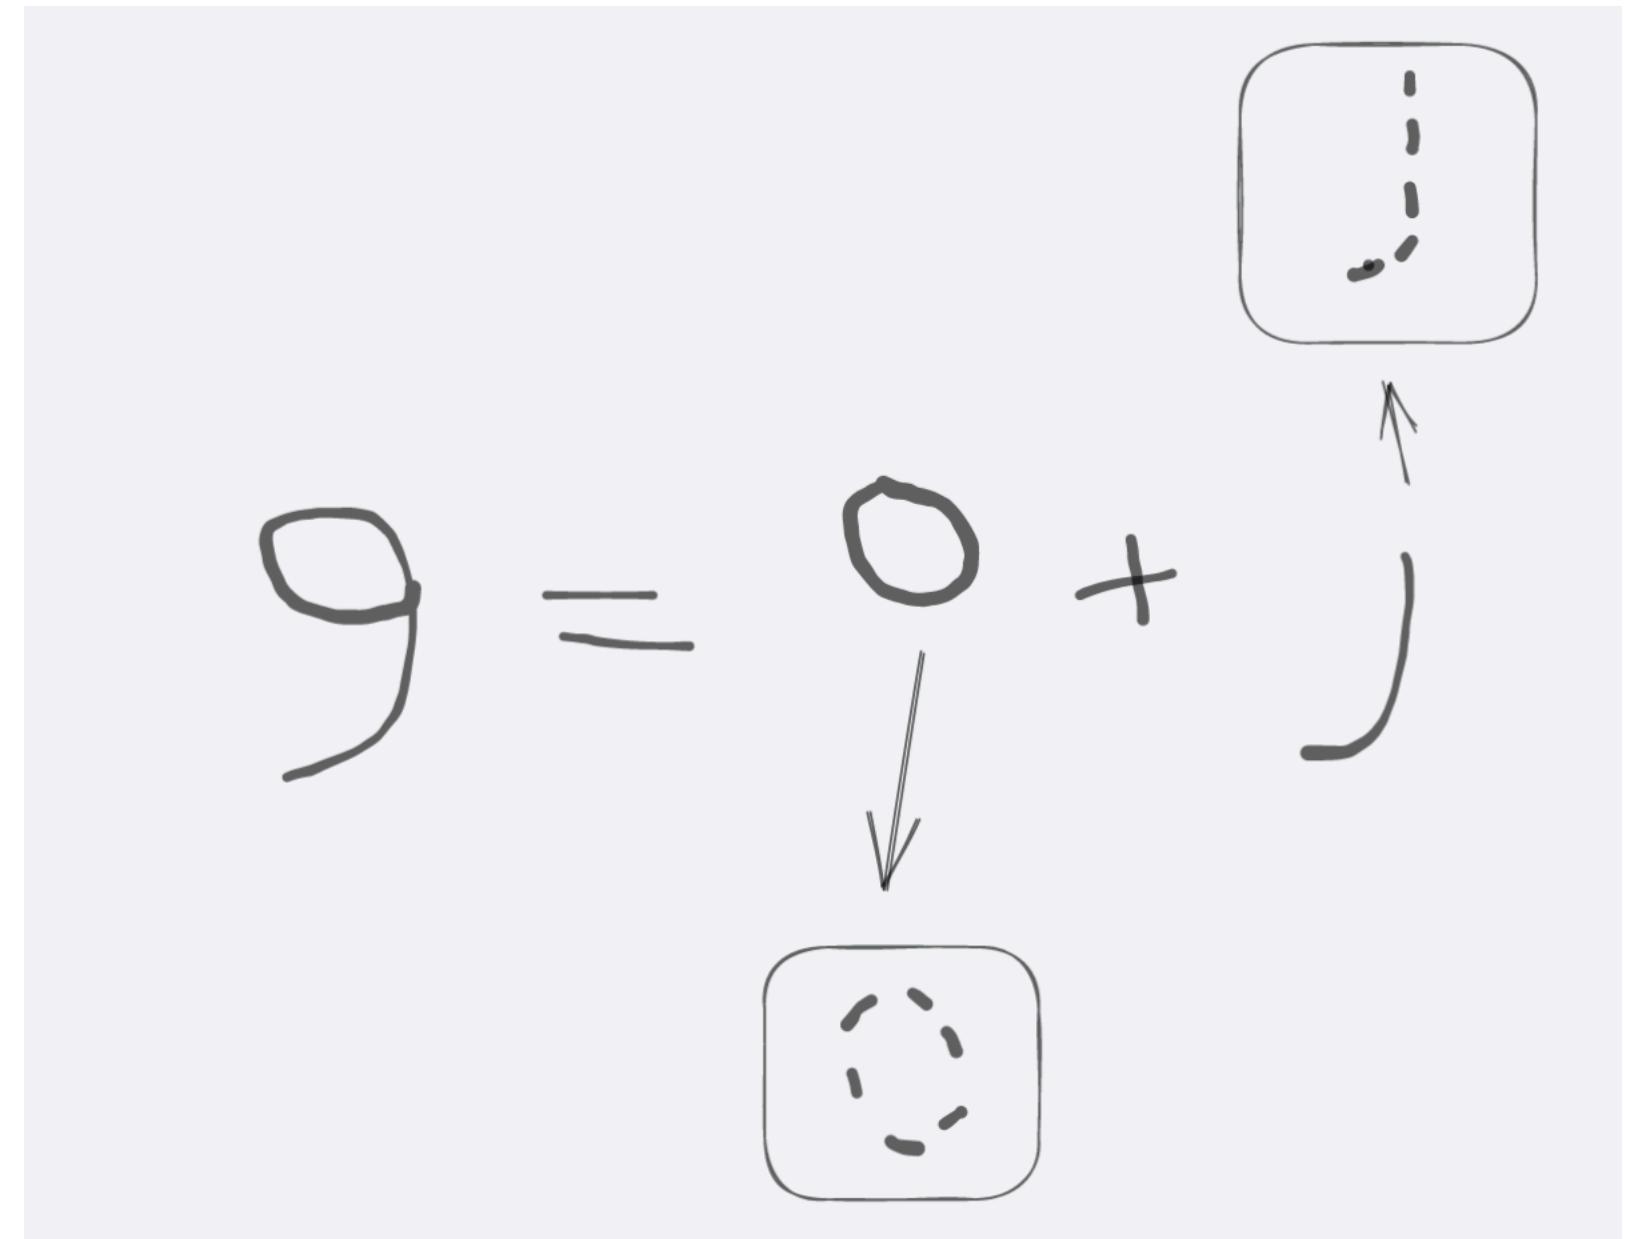


Recognition → ‘Re’ + ‘cog’ + ‘ni’ + ‘tion’

Neurones are mathematical function that collects & classifies information according to special architecture



Each Neurone is a perceptron, and is similar to
Multiple linear regression

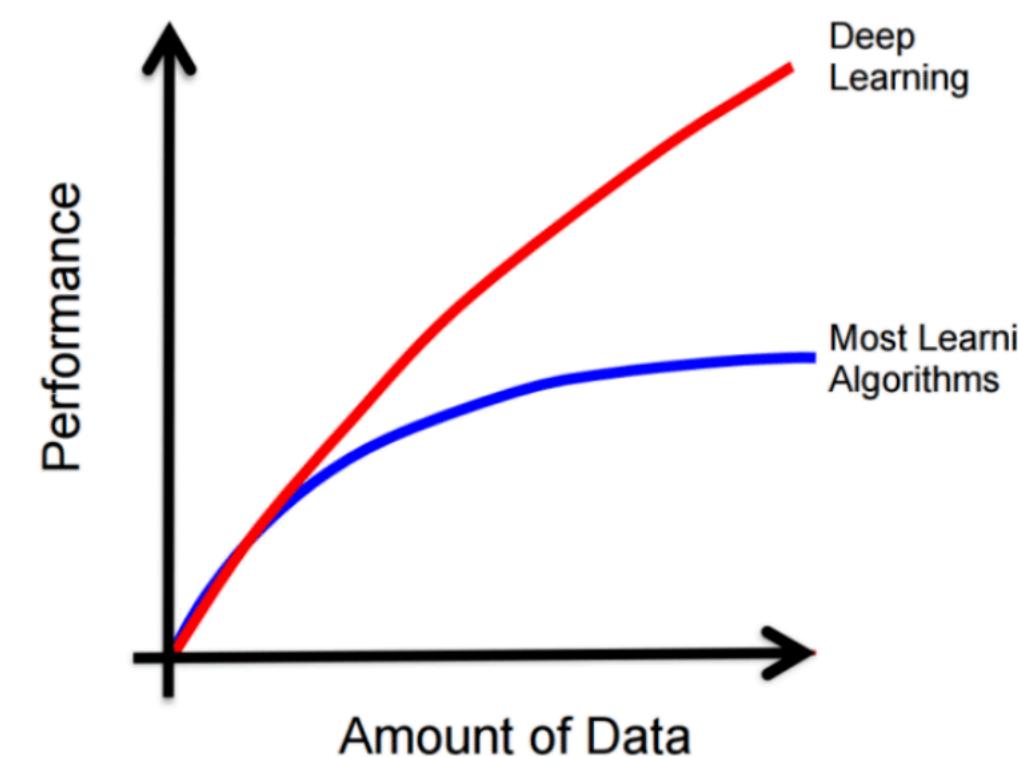
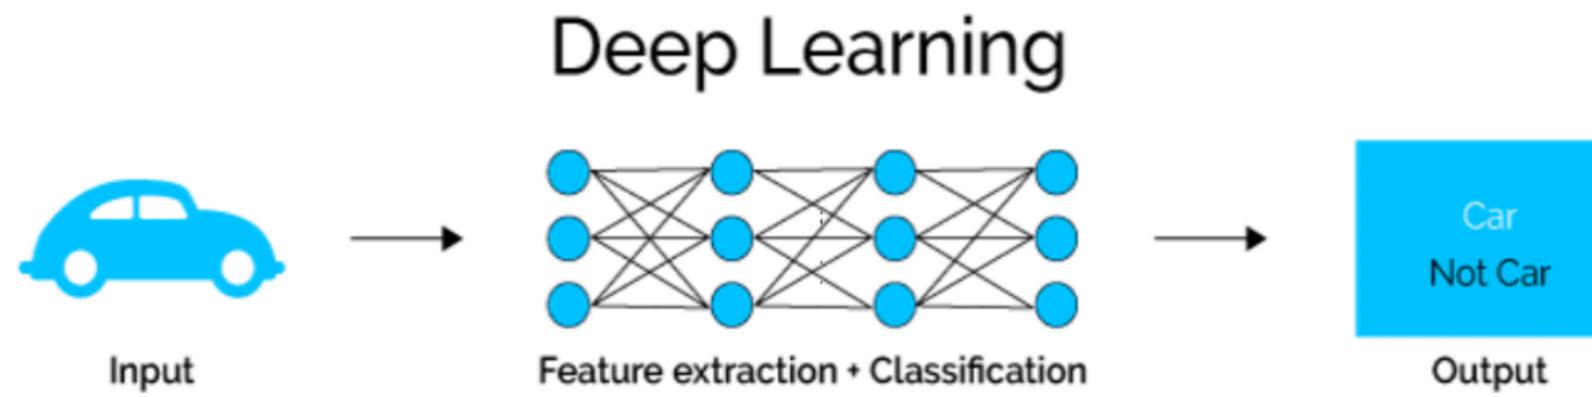
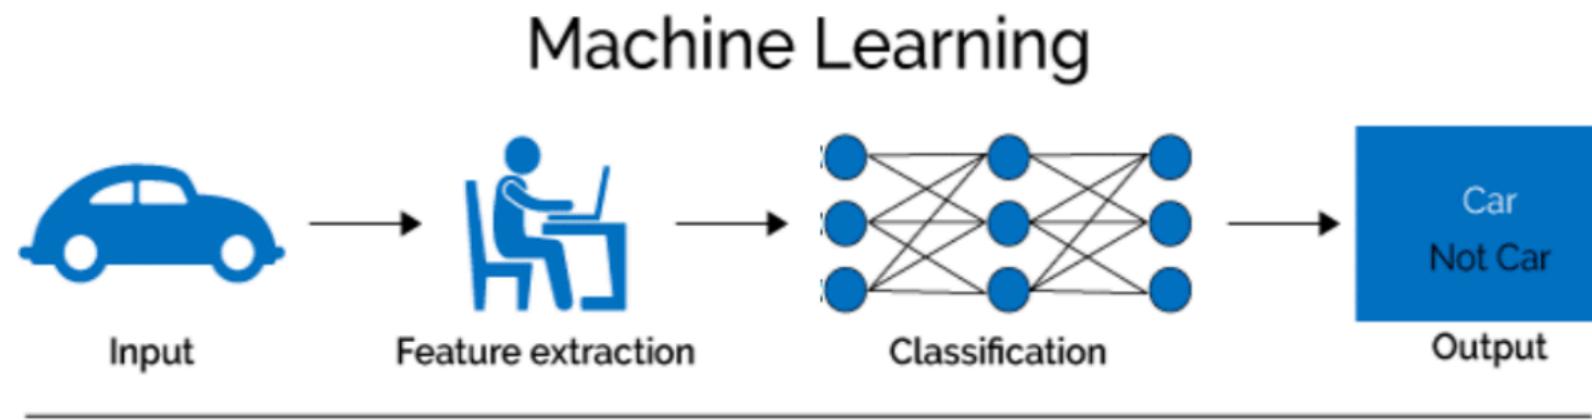


Types of Neural Networks

1. Feed-forward Neural Network
2. Recurrent Neural Network
3. Convolutional Neural Network
4. Deconvolutional Neural Network
5. Modular Neural Network

Deep Neural Network containing multiple layers, does its feature scaling itself.

There's no need to extract features from raw data (like sentiments in a sentence)



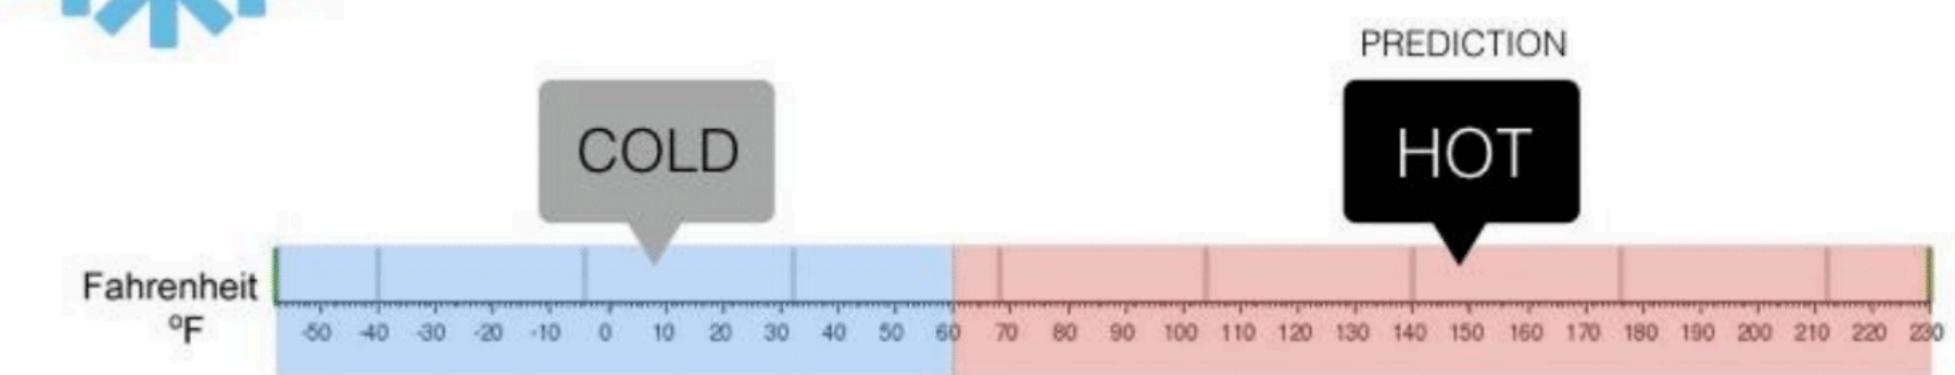
Regression

What is the temperature going to be tomorrow?



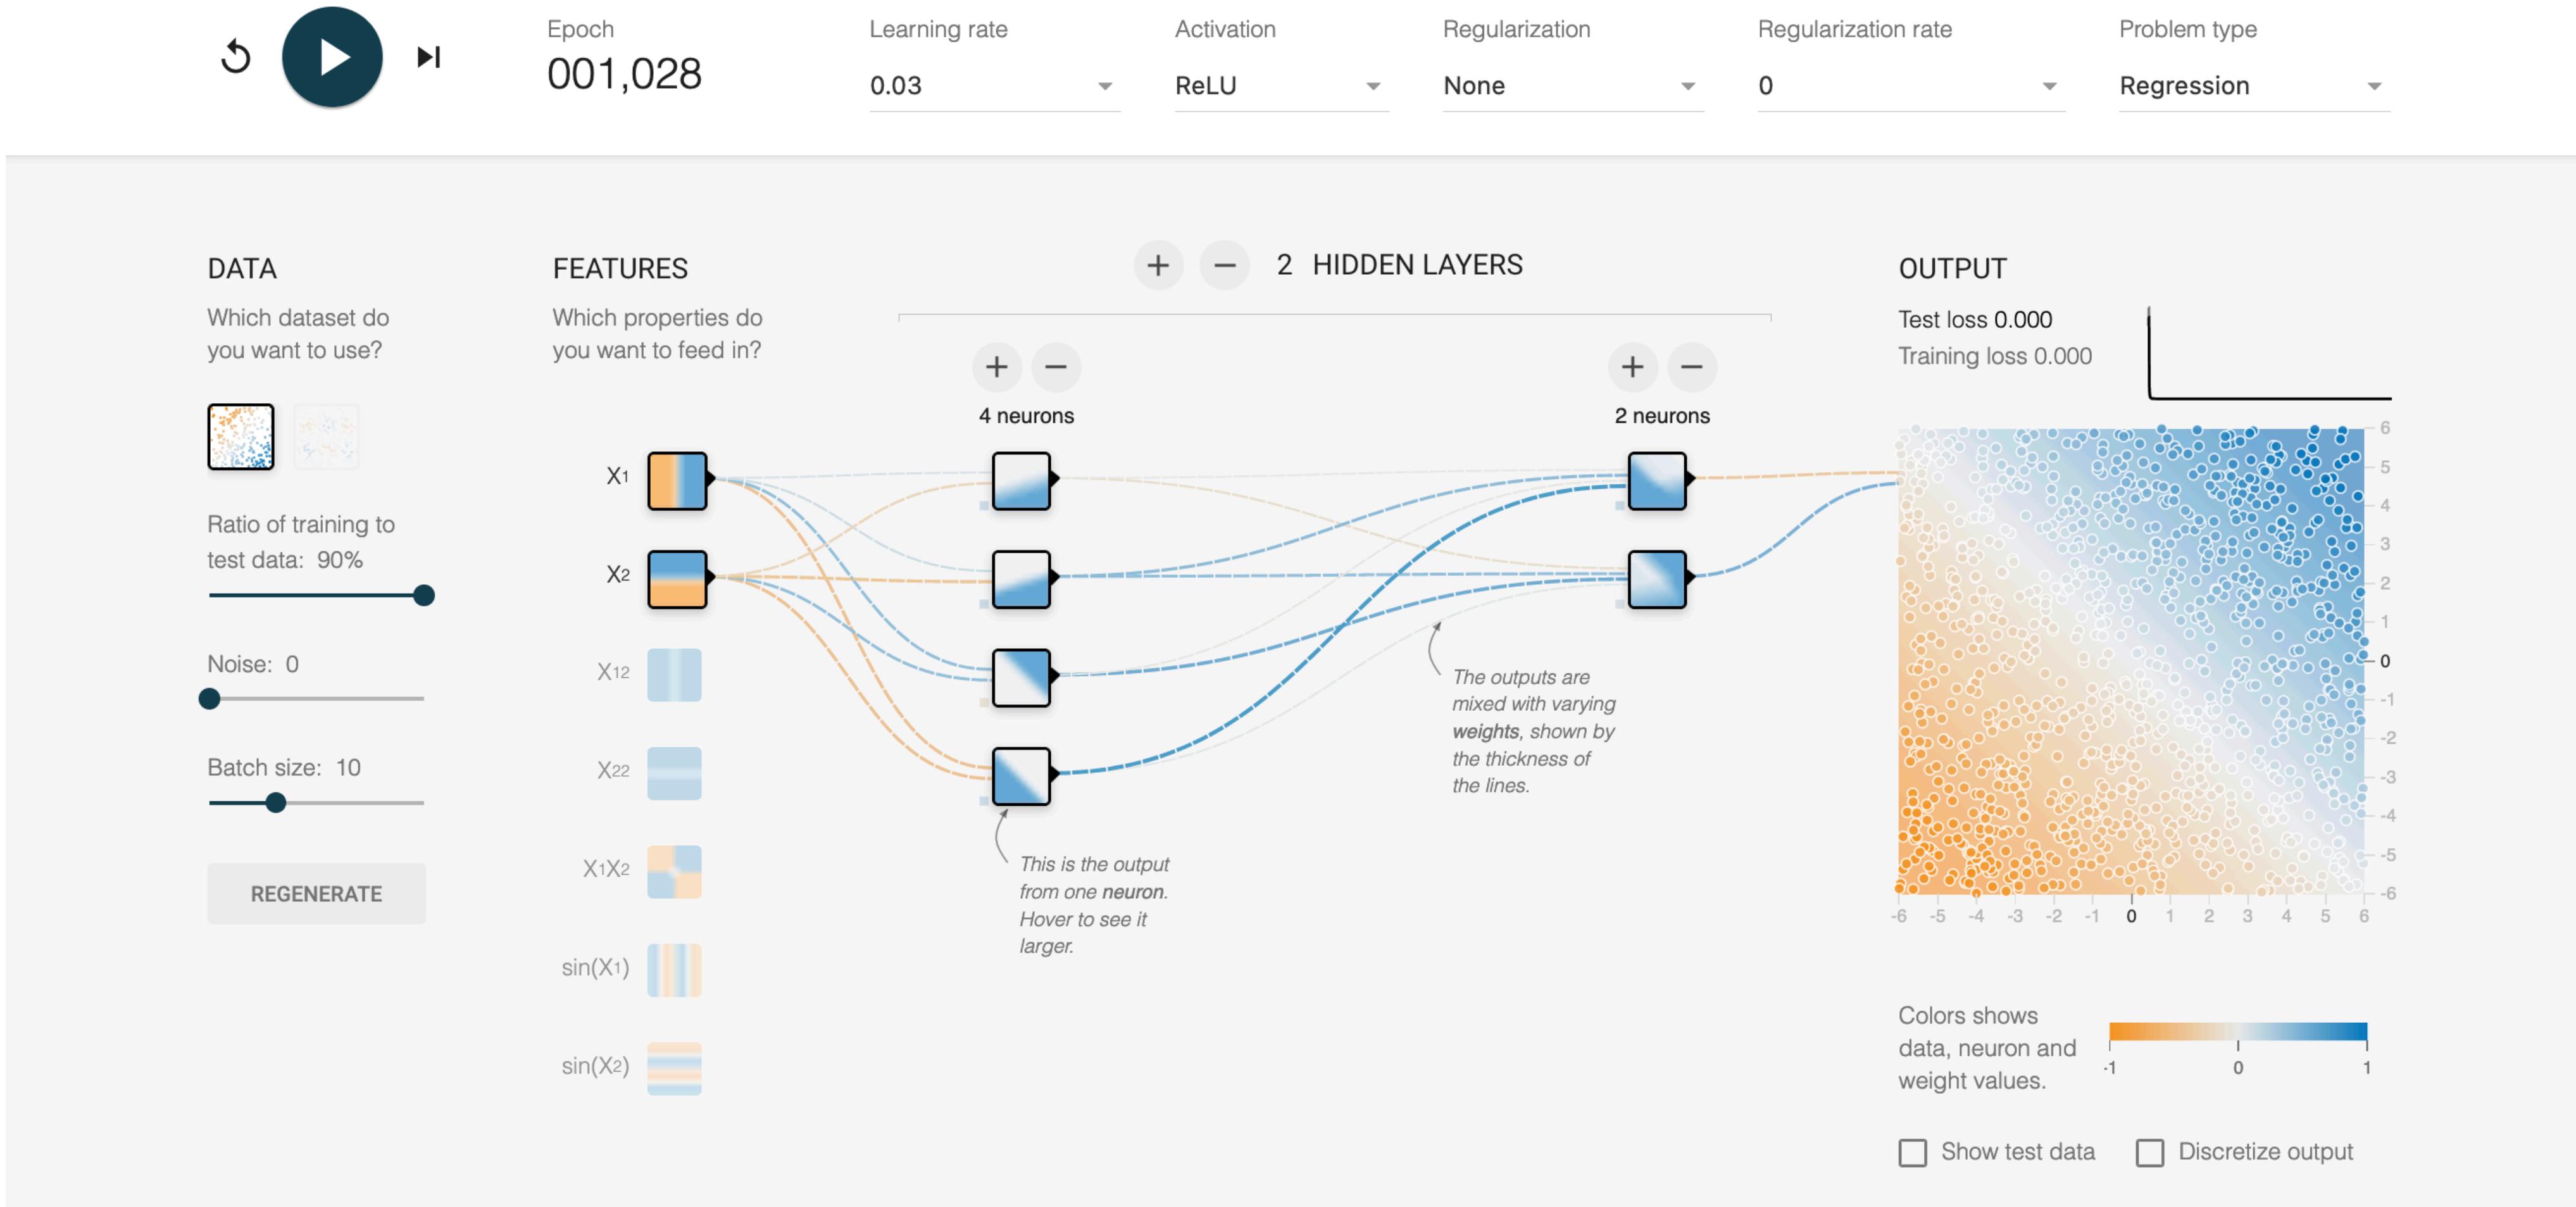
Classification

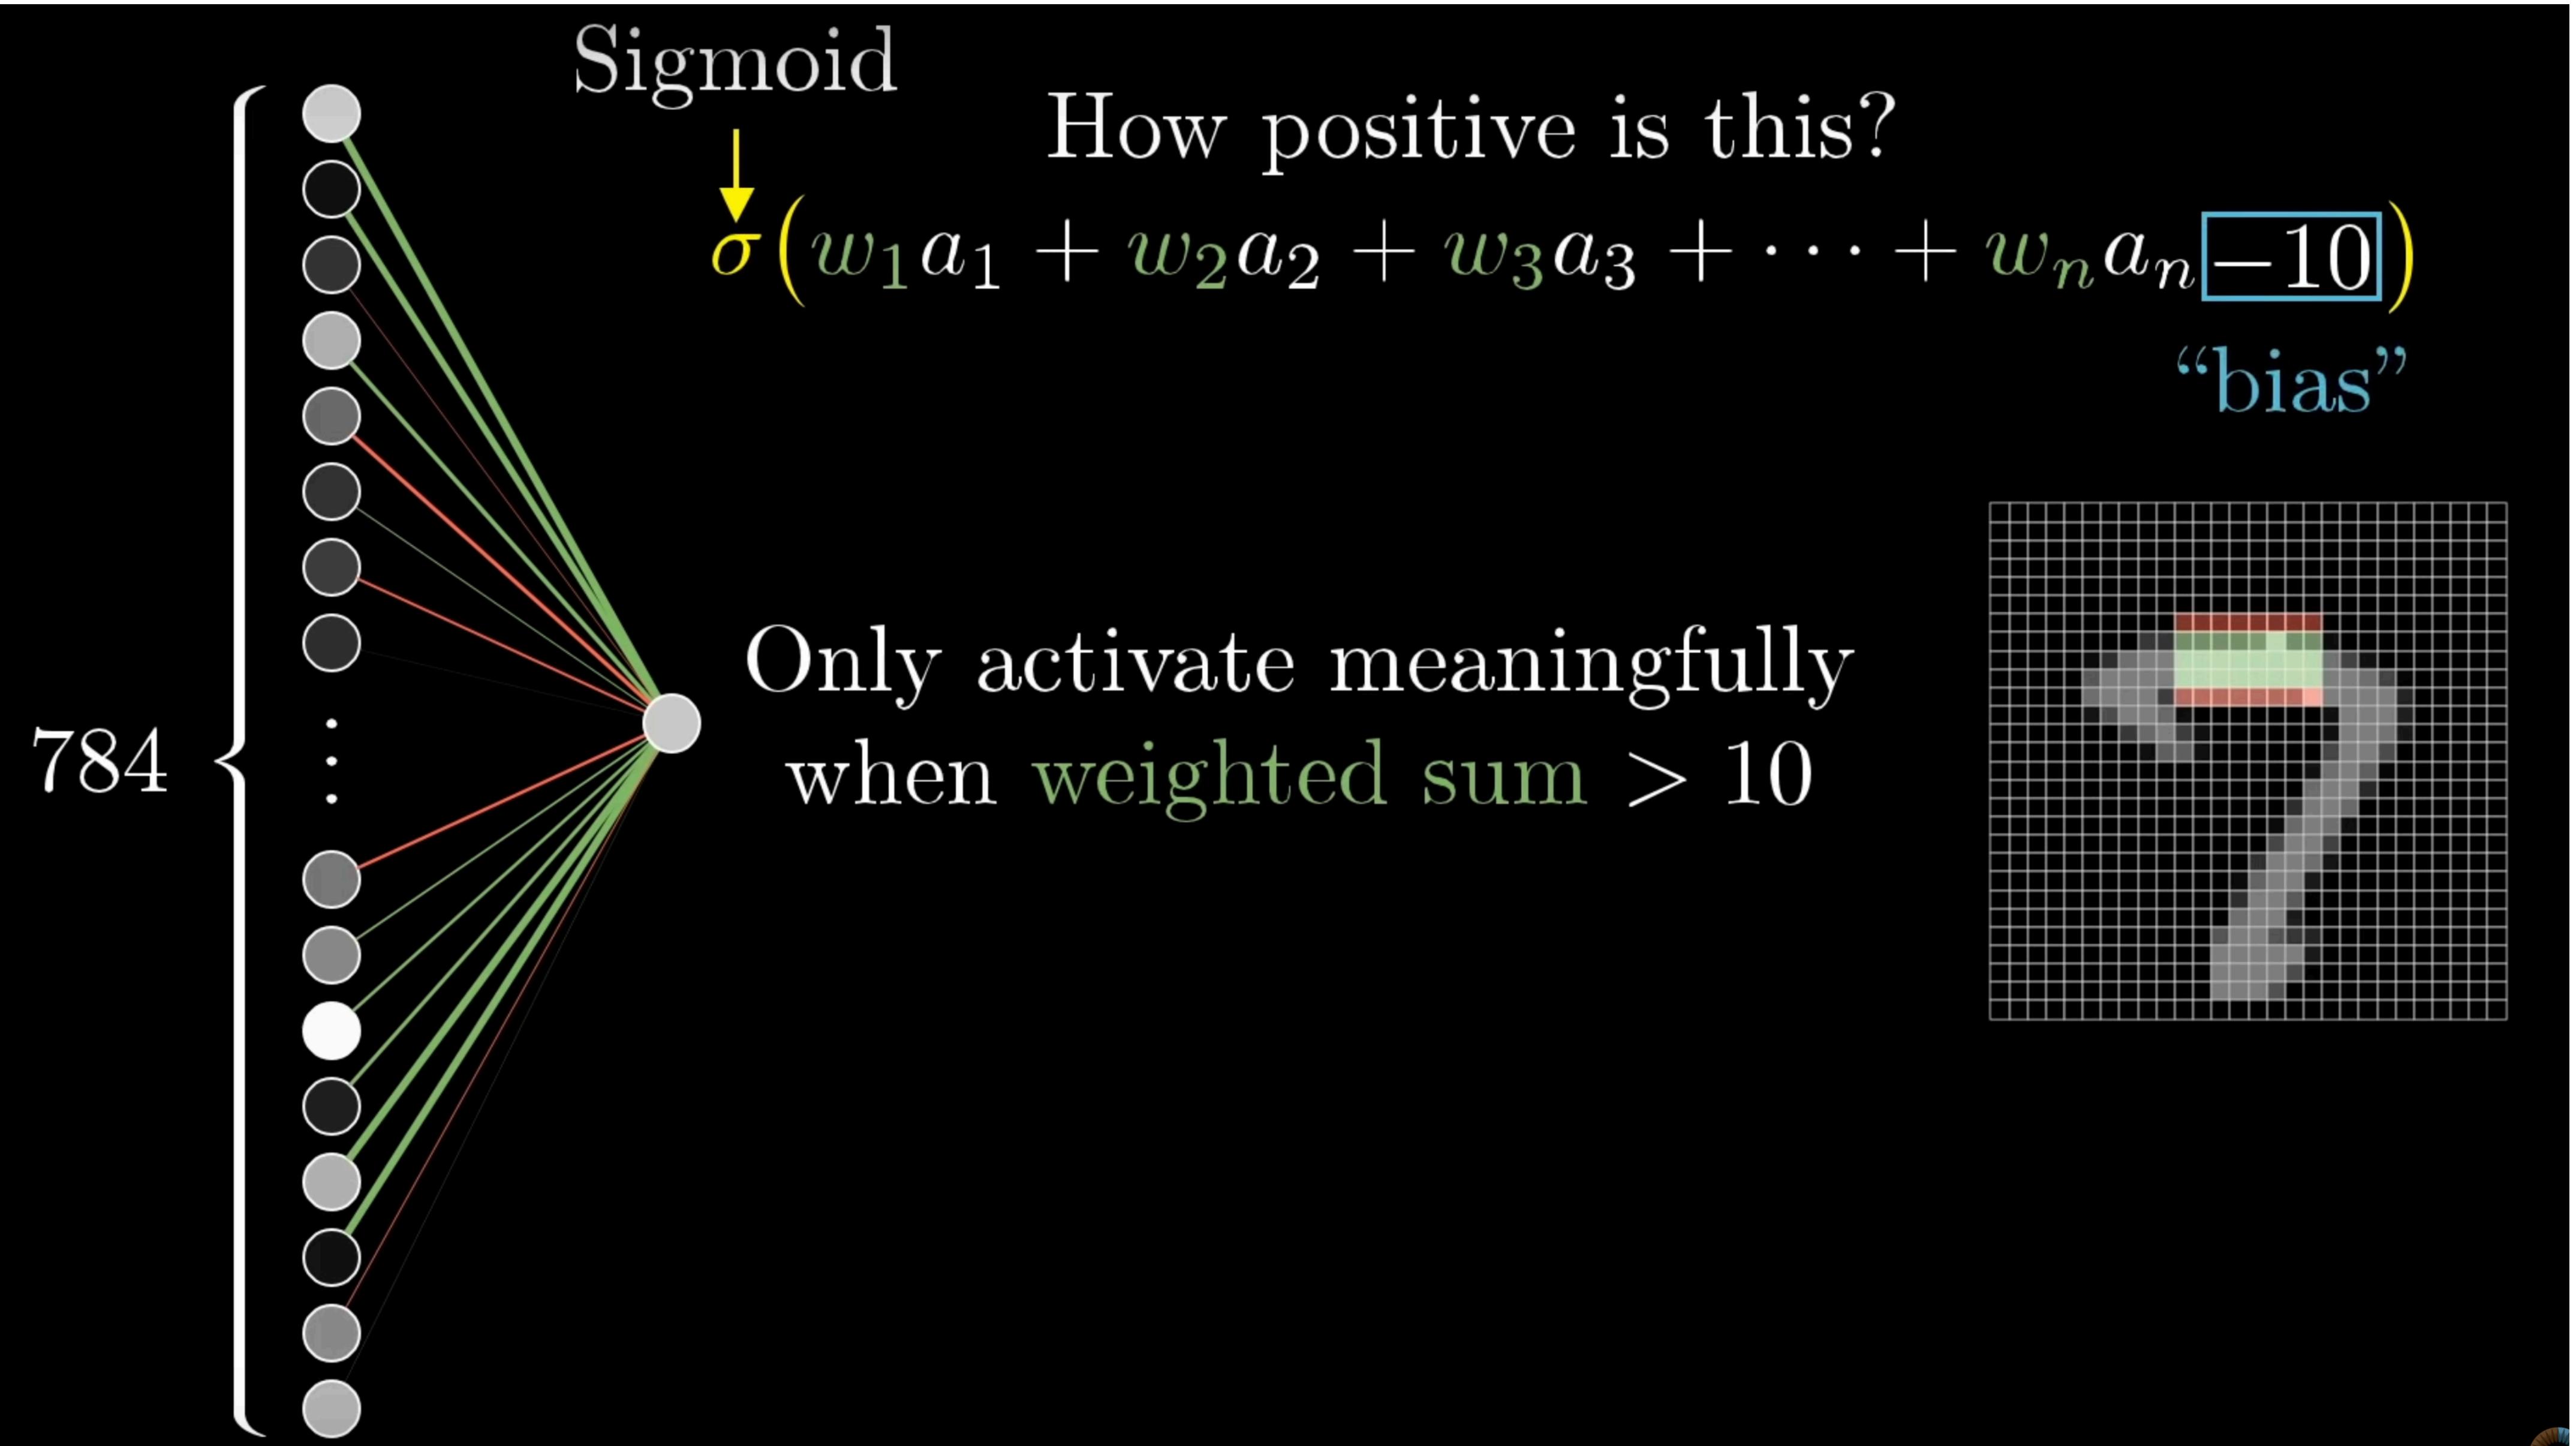
Will it be Cold or Hot tomorrow?



Tinker With a Neural Network Right Here in Your Browser.

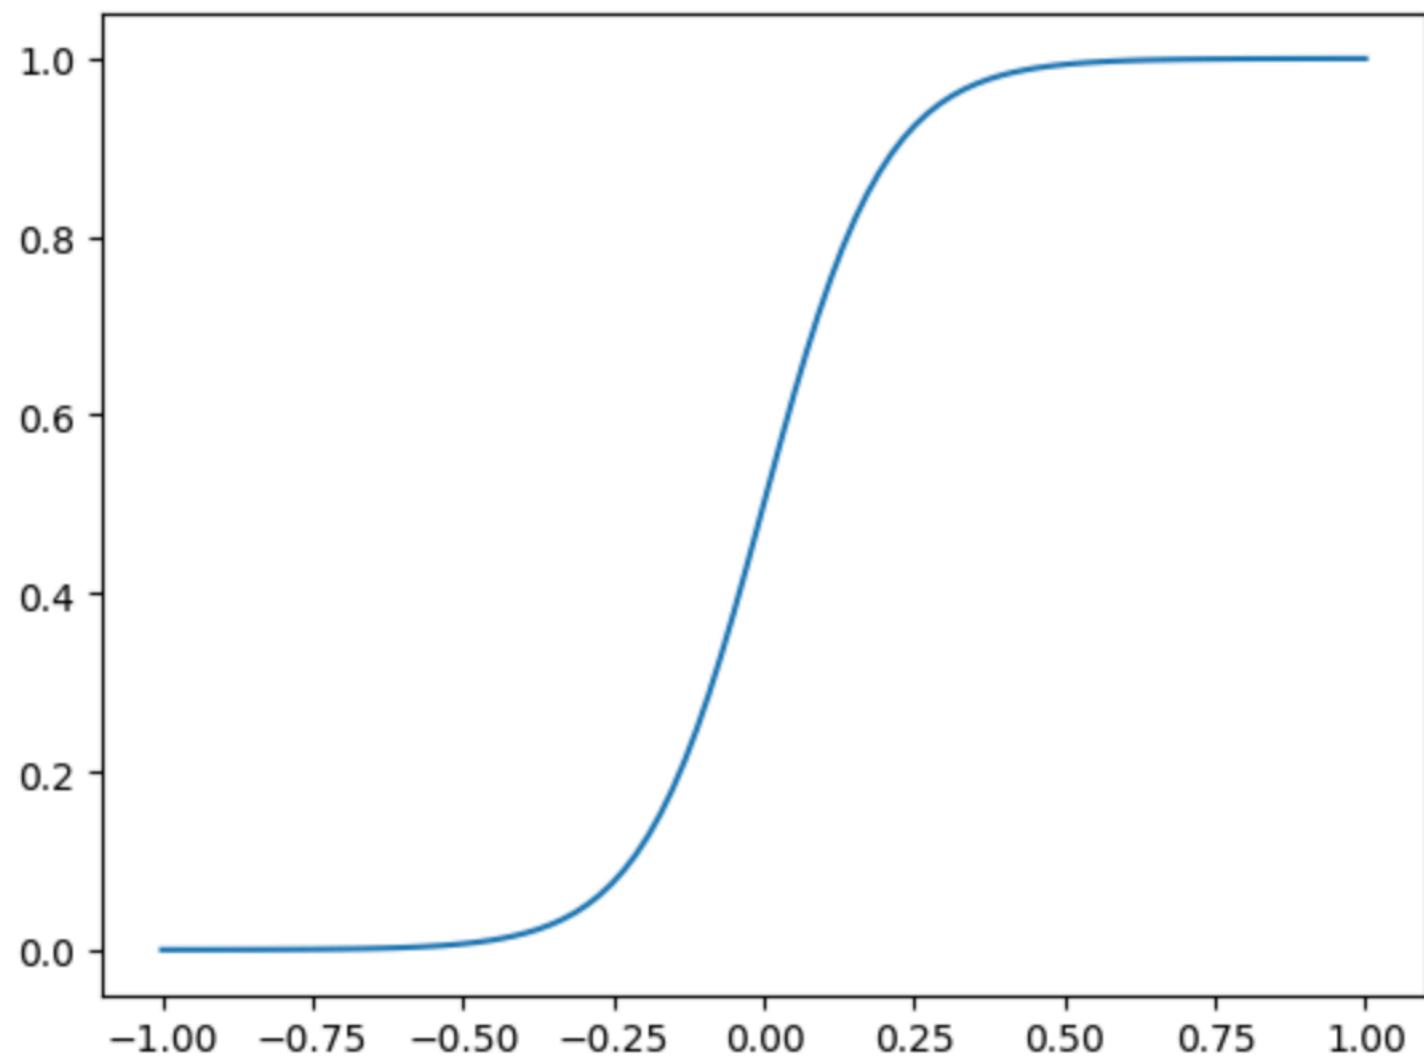
Don't Worry, You Can't Break It. We Promise.





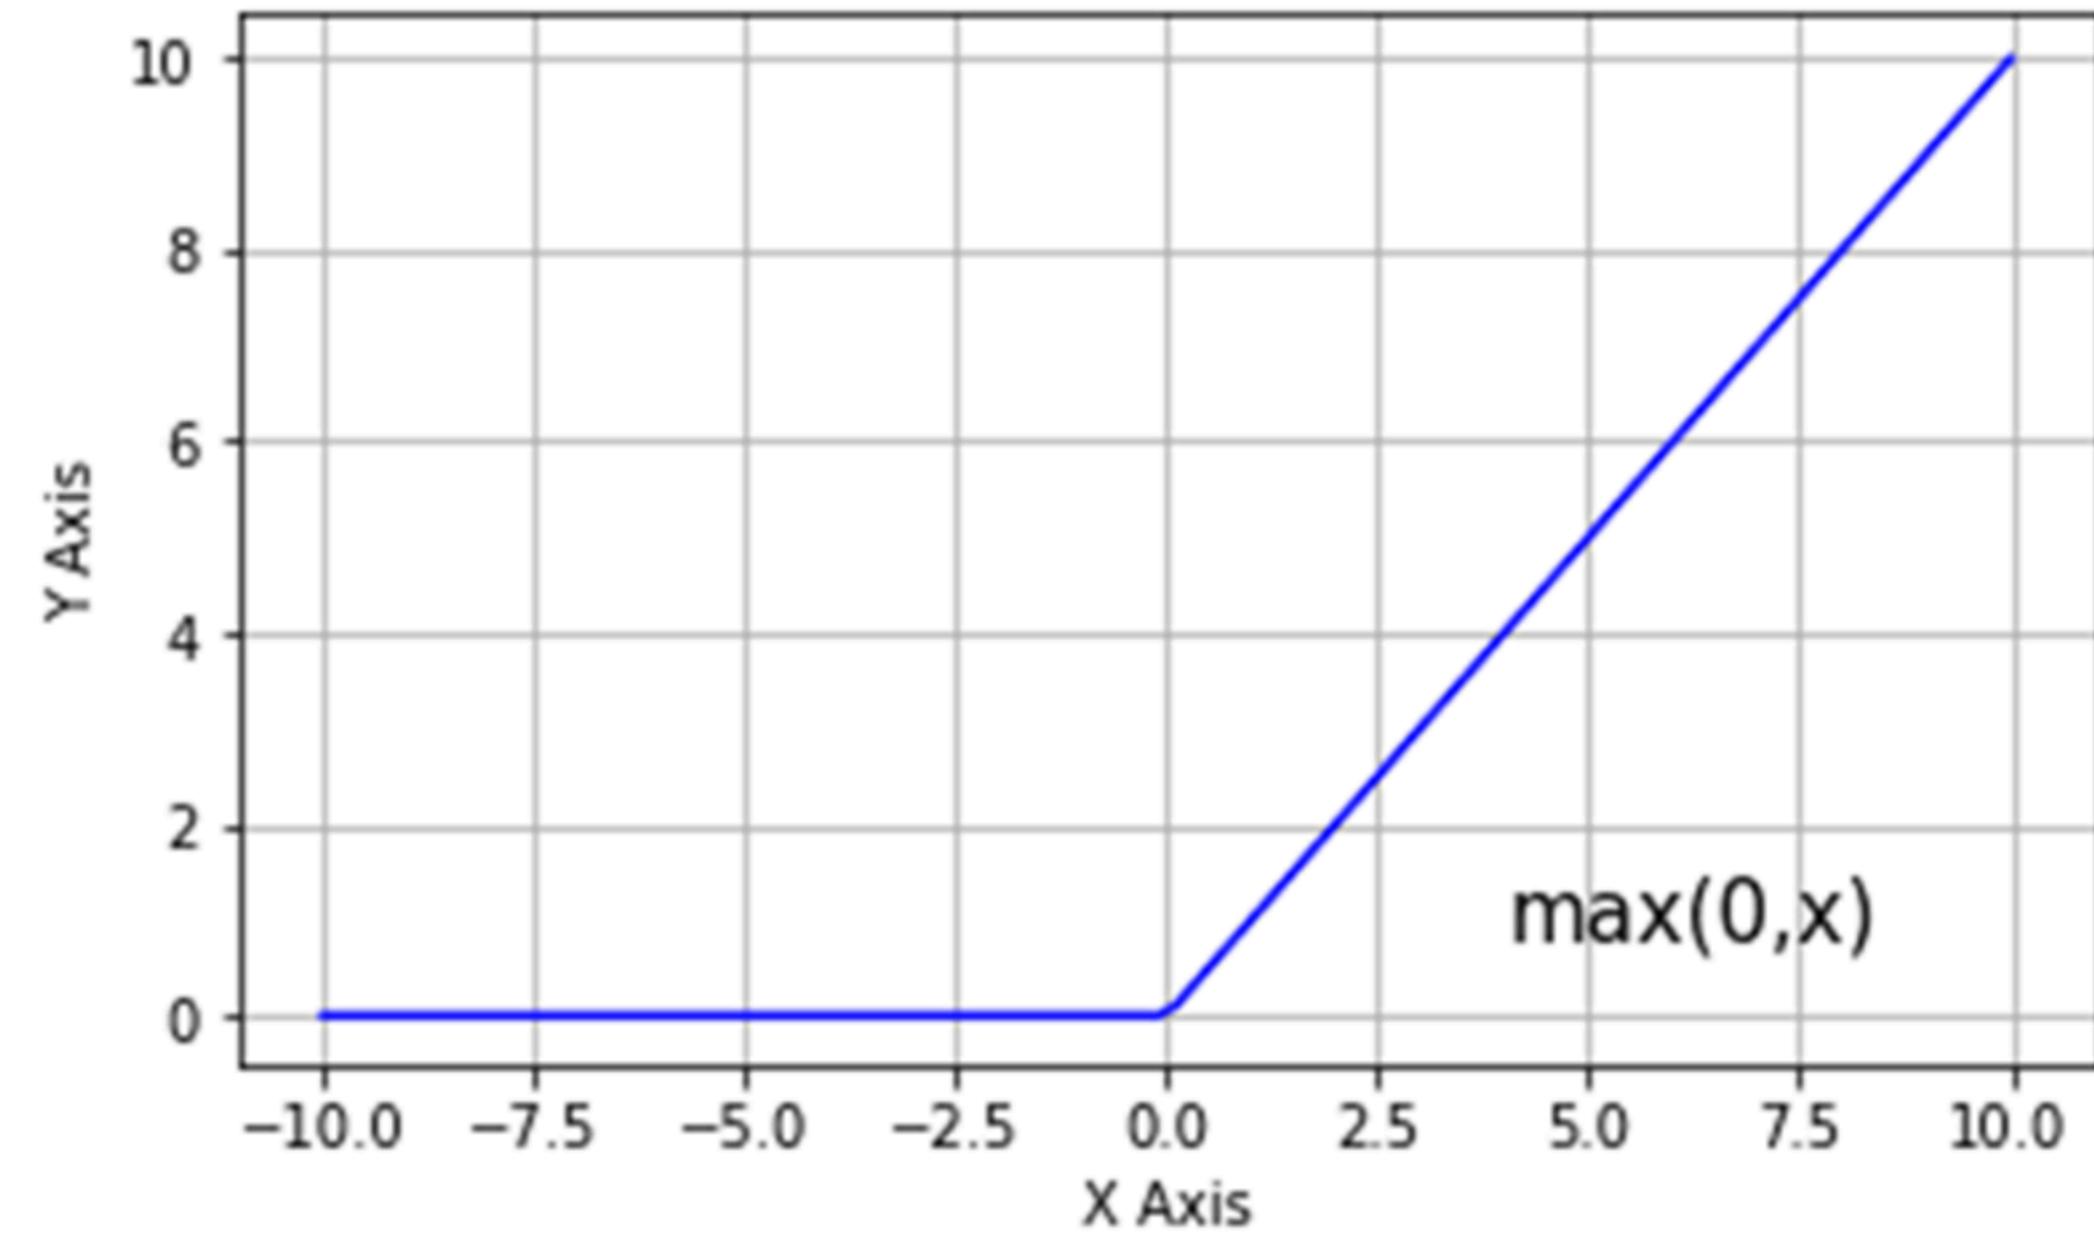
Activation Functions

Sigmoid Function



$$A = \frac{1}{1+e^{-x}}$$

ReLU Activation Function

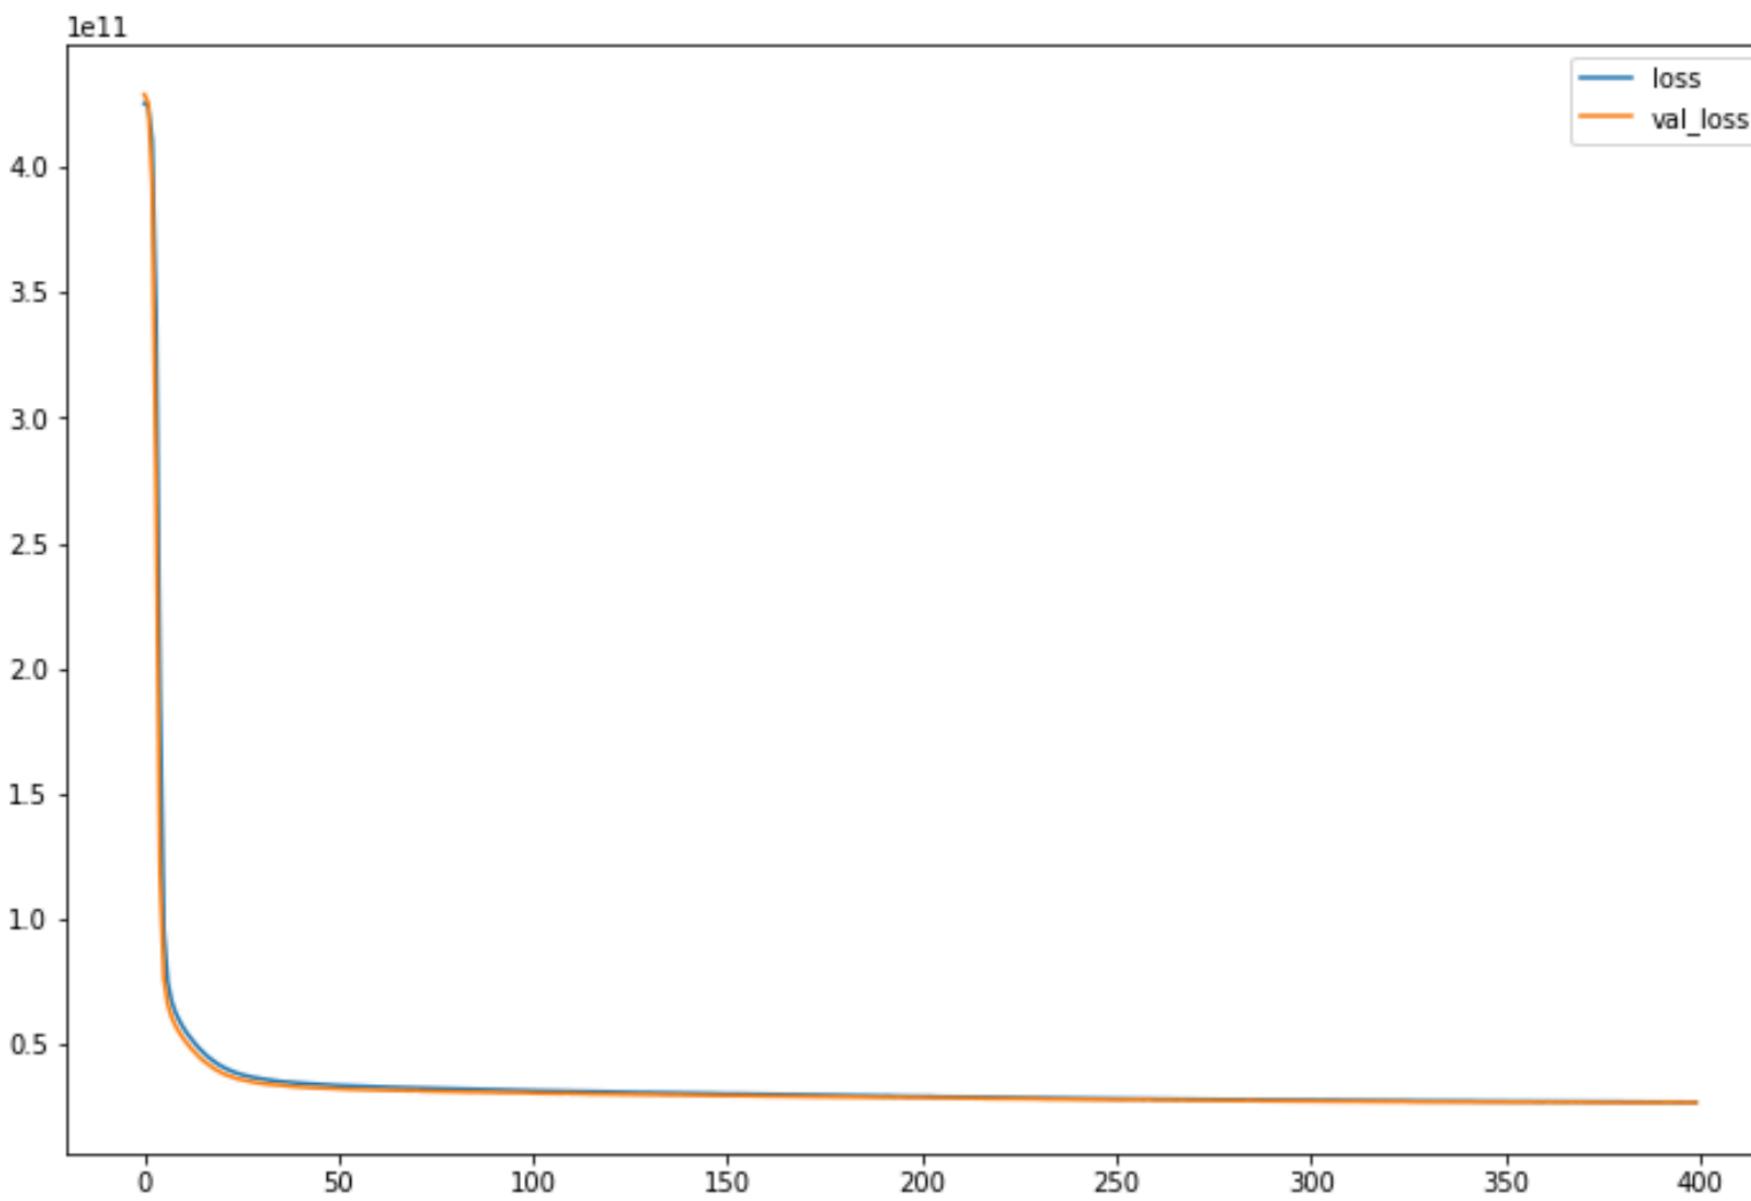


$$RELU(x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$$

The purpose of the activation function is to introduce non-linearity into the output of the neuron, allowing the neural network to learn complex patterns and relationships in the input data.

Implementation in Python

- Import Sequential, Dense, Activation & Adam (Adaptive Moment Estimation)
- Train the neural network
- Use fit() method to feed the training data into the network (x_train, y_train)
- Validate with the test datas (x_test, y_test)
- Declare batch size & epochs
- Plot the loss function with respect to training time
- Generate predictions
- Find accuracy & errors



The goal of training a neural network is to minimize the value of the loss function. This is done by adjusting the weights and biases of the network in a process called backpropagation.

```
156 157 158 159 160 161 162 163 164 165 166 167 168 169 170 171 172 173 174 175 176 177 178 179 180 181 182 183 184 185 186 187 188 189 190 191 192 193 194 195 196 197 198 199 200 201
```

Price_Prediction.py > ...

```
'''# Keras Regression'''

# Creating a Neural Network Model
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Activation
from tensorflow.keras.optimizers import Adam

# having 19 nueron is based on the number of available features

model = Sequential()

model.add(Dense(19,activation='relu'))
model.add(Dense(19,activation='relu'))
model.add(Dense(19,activation='relu'))
model.add(Dense(19,activation='relu'))
model.add(Dense(1))

model.compile(optimizer='adam', loss='mse')

'''## Training the mdoel'''

model.fit(x=X_train,y=y_train,
           validation_data=(X_test,y_test),
           batch_size=128,epochs=400)

model.summary()

loss_df = pd.DataFrame(model.history.history)
loss_df.plot(figsize=(12,8))

'''### Evaluation on Test Data'''

y_pred = model.predict(X_test)

# evaluation metrics
# explained variance score: best possible score is 1 and lower values are worse
from sklearn import metrics

print('Mean Absolute Error: {:.2f}'.format(metrics.mean_absolute_error(y_test, y_pred)))
print('Mean Squared Error: {:.2f}'.format(metrics.mean_squared_error(y_test, y_pred)))
print('Root Mean Squared Error: {:.2f}'.format(np.sqrt(metrics.mean_squared_error(y_test, y_pred))))
print('Variance score is: {:.2f}'.format(metrics.explained_variance_score(y_test,y_pred)))

# we are off about 20% (comparing mean absolut error and mean of price)
Data['price'].mean()
```

Accuracy

**Multiple Regression Model
(70.4%)**

**Neural Network using Keras
(80.51%)**

```
print('Model: Keras Regression\n')

print('Mean Absolute Error(MAE): {:.2f}'.format(metrics.mean_absolute_error(y_test, y_pred)))
print('Mean Squared Error(MSE): {:.2f}'.format(metrics.mean_squared_error(y_test, y_pred)))
print('Root Mean Squared Error(RMSE): {:.2f}'.format(np.sqrt(metrics.mean_squared_error(y_test, y_pred))))
print('Variance score: {:.2f}\n'.format(metrics.explained_variance_score(y_test,y_pred)*100))
print('*****\n')
print('Model: Multiple Linear Regression\n')
print('Mean Absolute Error(MAE): {:.2f}'.format(metrics.mean_absolute_error(y_test, y_predd)))
print('Mean Squared Error(MSE): {:.2f}'.format(metrics.mean_squared_error(y_test, y_predd)))
print('Root Mean Squared Error(RMSE): {:.2f}'.format(np.sqrt(metrics.mean_squared_error(y_test, y_predd))))
print('Variance score: {:.2f}'.format(metrics.explained_variance_score(y_test,y_predd)*100))
```

Model: Keras Regression

Mean Absolute Error(MAE): 101572.44
Mean Squared Error(MSE): 26737735523.03
Root Mean Squared Error(RMSE): 163516.77
Variance score: 80.51

Model: Multiple Linear Regression

Mean Absolute Error(MAE): 125933.74
Mean Squared Error(MSE): 40601153229.62
Root Mean Squared Error(RMSE): 201497.28
Variance score: 70.40

Future Aspects

- Incorporating additional data sources, such as economic indicators or news articles, to improve the accuracy of the predictions.
- Experimenting with different neural network architectures, such as recurrent neural networks or transformers, to see if they perform better than the current architecture.
- Developing an online or real-time version of the improved model.
- Applying the model to different markets or assets, such as stocks, commodities, or currencies, to see if the model is applicable to other prediction tasks.



References

1. "Deep Learning Basics: Introduction and Overview". MIT course 6.S094. Lex Fridman. <https://youtu.be/O5xeyoRL95U>
2. Quang Truong, Minh Nguyen, Hy Dang, Bo Mei, "Housing Price Prediction via Improved Machine Learning Techniques", Procedia Computer Science, Volume 174, 2020, Pages 433-442, ISSN 1877-0509, <https://doi.org/10.1016/j.procs.2020.06.111>
3. Zulkifley, Nor & Rahman, Shuzlina & Nor Hasbiah, Ubaidullah & Ibrahim, Ismail. (2020). "House Price Prediction using a Machine Learning Model: A Survey of Literature". International Journal of Modern Education and Computer Science. 12. 46-54, <https://doi.org/10.5815/ijmecs.2020.06.04>
4. Qingqi Zhang, "Housing Price Prediction Based on Multiple Linear Regression", Scientific Programming, vol. 2021, Article ID 7678931, 9 pages, 2021. <https://doi.org/10.1155/2021/7678931>
5. "Neural Regression using Keras". <https://visualstudiomagazine.com/articles/2018/07/23/neural-regression-using-keras.aspx>
6. "Regression with Keras". Pyimagesearch. <https://pyimagesearch.com/2019/01/21/regression-with-keras/>
7. "House Sales in King County, USA". Kaggle. <https://www.kaggle.com/datasets/harlfoxem/housesalesprediction>
8. "Multiple Linear Regression Model, A Quick Guide". Scribbr. <https://www.scribbr.com/statistics/multiple-linear-regression/>
9. "But what is a neural network? | Chapter 1, Deep learning". 3Blue1Brown. <https://youtu.be/aircAruvnKk>
10. "Matplotlib References". <https://matplotlib.org/stable/api/index.html>
11. "Seaborn: Statistical Data Visualisation". <https://seaborn.pydata.org/>
12. "A Neural Network Playground". Tensorflow. <https://playground.tensorflow.org/>

Thank You